

Assignment 1 Deep Learning and Sensor fusion 2026

Multi class Classification

Necessary import here

```
In [1]: import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import DataLoader, Dataset
from torch.utils.data import Dataset
import torch.nn.functional as F
import matplotlib.pyplot as plt
from PIL import Image
from torchvision import transforms
import os
import numpy as np
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay
from sklearn.model_selection import train_test_split

print(f"[INFO] Torch infos: {torch.__version__}")

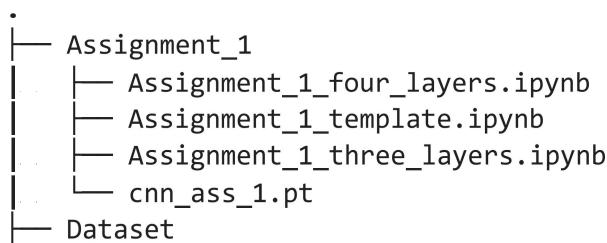
DATASET_PATH = os.path.join("../", "Dataset", "BoneBreakClassification")
```

[INFO] Torch infos: 2.7.1+cu118

```
In [2]: # some helpers
def load_data(dataset_path, class_names, type_of_data_to_load):
    # for each class name
    images, labels = list(), list()
    for el in class_names.keys():
        print(f"[WALKING] Walking into {el}")
        for images_name in os.listdir(os.path.join(dataset_path, el, type_of_data_t
            full_image_path = os.path.join(dataset_path, el, type_of_data_to_load,
            images.append(full_image_path)
            labels.append(class_names[el])
    return images, labels
```

set Dataset path

The path you see in the code point to my folder



```

| . └─ BoneBreakClassification

```

If you put the data in a different place change it.

```
In [3]: DATASET_PATH = os.path.join("../", "Dataset", "BoneBreakClassification")
```

(1) Load dataset

```

In [4]: class_names = dict()

temp = os.listdir(DATASET_PATH)
temp = sorted(temp)

# prep up list and class number together
for i in range(len(temp)):
    class_names[temp[i]] = i

data_to_split, labels_to_split = load_data(dataset_path=DATASET_PATH,
                                             class_names=class_names,
                                             type_of_data_to_load="Train")

print(f"[DATA TO SPLIT] Data loaded complete {len(data_to_split)} Labeles are {len(labels_to_split)}")

# Load test data
test_data_handler, test_labels_handler = load_data(dataset_path=DATASET_PATH,
                                                     class_names=class_names,
                                                     type_of_data_to_load="Test")

print(f"[TRAIN DATA LOADED] Data loaded complete {len(test_data_handler)} Labeles are {len(test_labels_handler)}")

```

```

[WALKING] Walking into Avulsion fracture
[WALKING] Walking into Comminuted fracture
[WALKING] Walking into Fracture Dislocation
[WALKING] Walking into Greenstick fracture
[WALKING] Walking into Hairline Fracture
[WALKING] Walking into Impacted fracture
[WALKING] Walking into Longitudinal fracture
[WALKING] Walking into Oblique fracture
[WALKING] Walking into Pathological fracture
[WALKING] Walking into Spiral Fracture
[DATA TO SPLIT] Data loaded complete 904 Labeles are 904
[WALKING] Walking into Avulsion fracture
[WALKING] Walking into Comminuted fracture
[WALKING] Walking into Fracture Dislocation
[WALKING] Walking into Greenstick fracture
[WALKING] Walking into Hairline Fracture
[WALKING] Walking into Impacted fracture
[WALKING] Walking into Longitudinal fracture
[WALKING] Walking into Oblique fracture
[WALKING] Walking into Pathological fracture
[WALKING] Walking into Spiral Fracture
[TRAIN DATA LOADED] Data loaded complete 131 Labeles are 131

```

(2) leave this cell because we need to create a val dataset and it is not there as default

```
In [5]: # here you can split validation in stead of test so you can preserve intacted the t
train_data_handler, validation_data_handler, train_labels_handler, validation_label
print(f"[TRAIN DATA] The train data is {train_data_handler.shape} and its labels a
print(f"[VALIDATION DATA] The validation data is {validation_data_handler.shape} a
# my data Loader need List yours? adapt it to your code
train_data_handler = list(train_data_handler)
train_labels_handler = list(train_labels_handler)

validation_data_handler = list(validation_data_handler)
validation_labels_handler = list(validation_labels_handler)

[TRAIN DATA] The train data is (723,) and its labels are (723,)
[VALIDATION DATA] The validation data is (181,) and its labels are (181,)
```

(3) Here are two transforms one for training and one for evaluation

train transform need data augmentation:

- add random flip
- add random rotation
- add color jitter

both need:

- resize
- toTensor handler
- Normalize (use the provided MEAN and STD)

Note that the evaluation transform does not need data augmentation.

```
In [6]: TARGET_H = 224
TARGET_W = 224
IMAGE_MEAN = [0.485, 0.456, 0.406]
IMAGE_STD = [0.229, 0.224, 0.225]

# create transform
train_transform = transforms.Compose([
    transforms.Resize((TARGET_H, TARGET_W)),
    transforms.RandomHorizontalFlip(),
    transforms.RandomVerticalFlip(),
    transforms.RandomRotation(degrees=15),
    transforms.ColorJitter(brightness=0.1, contrast=0.1),
```

```

        transforms.ToTensor(),
        transforms.Normalize(IMAGE_MEAN, IMAGE_STD),
    ])

evaluation_transform = transforms.Compose([
    transforms.Resize((TARGET_H, TARGET_W)),
    transforms.ToTensor(),
    transforms.Normalize(IMAGE_MEAN, IMAGE_STD),
])

```

```

In [7]: class MyDataset(Dataset):
    def __init__(self, list_of_pathes, list_of_classes, transform=None) -> None:
        super().__init__()
        self.list_of_pathes = list_of_pathes
        self.list_of_classes = list_of_classes
        self.transform = transform

    def __len__(self):
        return len(self.list_of_pathes)

    def __getitem__(self, index):
        # select the image at index
        path = self.list_of_pathes[index]
        # select the label at index
        label = self.list_of_classes[index]
        # open the image Advice use Image package
        image = Image.open(path).convert("RGB")
        # apply transform
        if self.transform:
            image = self.transform(image)
        # return image and label
        return image, label

```

(4) Create data loder from the MyDataset class

```

In [8]: train_dataset = MyDataset(list_of_pathes=train_data_handler,
                               list_of_classes=train_labels_handler,
                               transform=train_transform)

validation_dataset = MyDataset(list_of_pathes=validation_data_handler,
                               list_of_classes=validation_labels_handler,
                               transform=evaluation_transform)

test_dataset = MyDataset(list_of_pathes=test_data_handler,
                        list_of_classes=test_labels_handler,
                        transform=evaluation_transform)

# put those here
train_loader = DataLoader(train_dataset, batch_size=64, shuffle=True, num_workers=0
validation_loader = DataLoader(validation_dataset, batch_size=128, shuffle=False, n
test_loader = DataLoader(test_dataset, batch_size=128, shuffle=False, num_workers=0

```

Instance of the CNN

Use the Sequential API it is faster

- This cnn has two hidden layers (conv2d -> proper activation -> max pool 2d) one has 16 units and the second has 32 units
- Add a flatten layer
- Classifier need 128 Linear unit proper activation and set dropdown at 0.5
- The last layer of the classifier how many units need?

```
In [9]: class LayeredCNN(nn.Module):
    def __init__(self,
                 layers_list: list[int],
                 dense_units: list[int],
                 numbers_of_output_classes: int = 10,
                 kernel_size: int = 3,
                 padding_size: int = 1,
                 pooling_size: int = 2,
                 input_channels: int = 3,
                 dropout_probabilities: float = 0.5):
        super().__init__()

        layers = []
        input_units = input_channels

        for output_units in layers_list:

            layers += [
                nn.Conv2d(input_units, output_units, kernel_size=kernel_size, padding=padding_size),
                nn.ReLU(inplace=True),
                nn.MaxPool2d(kernel_size=pooling_size)
            ]

            input_units = output_units

        self.features = nn.Sequential(*layers)

        multi_layer_perceptron = []

        in_features = dense_units[0]

        for h in dense_units[1:]:
            multi_layer_perceptron += [
                nn.Linear(in_features, h),
                nn.ReLU(inplace=True),
                nn.Dropout(p=dropout_probabilities),
            ]

            in_features = h

        self.classifier = nn.Sequential(*multi_layer_perceptron)
        self.flatten_layer = nn.Flatten()
        self.final_classifier = nn.Linear(in_features, numbers_of_output_classes)

    def forward(self, x):
        x = self.features(x)
```

```

        x = self.flatten_layer(x)
        x = self.classifier(x)
        x = self.final_classifier(x)

    return x

```

(5) Add proper optimizer and loss function set training epochs to 50

```

In [10]: model = LayeredCNN([16,32], [100352, 128])

device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

model.to(device)

print(f"Model is running on: {device}")

criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(), lr=1e-3, weight_decay=1e-4)

```

Model is running on: cuda

```

In [11]: def train_one_epoch(model, loader, criterion, optimizer, device):
    model.train()
    running_loss = 0.0
    correct = 0
    total = 0

    for images, labels in loader:
        images, labels = images.to(device), labels.to(device)

        optimizer.zero_grad()
        outputs = model(images)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()

        # Track stats
        running_loss += loss.item() * images.size(0)
        predictions = outputs.argmax(dim=1)
        correct += (predictions == labels).sum().item()
        total += labels.size(0)

    return running_loss / total, correct / total

```

```

In [12]: @torch.no_grad()

def evaluate(model, loader, criterion, device):
    model.eval() # Turn off Dropout / BatchNorm
    running_loss = 0.0
    correct = 0
    total = 0

    for images, labels in loader:
        images, labels = images.to(device), labels.to(device)

```

```

outputs = model(images)
loss = criterion(outputs, labels)
running_loss += loss.item() * images.size(0)
predictions = outputs.argmax(dim=1)
correct += (predictions == labels).sum().item()
total += labels.size(0)

return running_loss / total, correct / total

```

(6) add here the train and evaluation loop collect loss and accuracy, we want to see curves here

```

In [13]: training_loss = []
validation_loss = []
training_accuracy = []
validation_accuracy = []

epochs = 50
best_validation_accuracy = 0.0

for epoch in range(1, epochs + 1):
    train_loss, train_accuracy = train_one_epoch(model = model, loader=train_loader,
                                                val_loss, val_accuracy = evaluate(model=model, loader=validation_loader, criterion=criterion)

    training_loss.append(train_loss)
    validation_loss.append(val_loss)
    training_accuracy.append(train_accuracy)
    validation_accuracy.append(val_accuracy)

    print(f"[TRAINING EPOCH {epoch} / {epochs}] TRAIN LOSS {train_loss: .4f}, TRAIN ACCURACY {train_accuracy: .4f}")
    print(f"[VALIDATION EPOCH {epoch} / {epochs}] VAL LOSS {val_loss: .4f}, VAL ACCURACY {val_accuracy: .4f}")

    if val_accuracy > best_validation_accuracy:
        best_validation_accuracy = val_accuracy
        torch.save(model.state_dict(), "cnn_bone_break_classification.pt")

print("TRAINING COMPLETE!!!")
print("BEST VALIDATION ACCURACY: ", best_validation_accuracy)

```

```
[TRAINING EPOCH 1 / 50] TRAIN LOSS 5.2374, TRAIN ACCURACY 0.098 | [VALIDATION EPOCH 1 / 50] VALIDATION LOSS 2.2866, VALIDATION ACCURACY 0.155
[TRAINING EPOCH 2 / 50] TRAIN LOSS 2.2952, TRAIN ACCURACY 0.136 | [VALIDATION EPOCH 2 / 50] VALIDATION LOSS 2.2912, VALIDATION ACCURACY 0.144
[TRAINING EPOCH 3 / 50] TRAIN LOSS 2.2804, TRAIN ACCURACY 0.134 | [VALIDATION EPOCH 3 / 50] VALIDATION LOSS 2.2833, VALIDATION ACCURACY 0.144
[TRAINING EPOCH 4 / 50] TRAIN LOSS 2.2704, TRAIN ACCURACY 0.138 | [VALIDATION EPOCH 4 / 50] VALIDATION LOSS 2.2806, VALIDATION ACCURACY 0.193
[TRAINING EPOCH 5 / 50] TRAIN LOSS 2.2762, TRAIN ACCURACY 0.137 | [VALIDATION EPOCH 5 / 50] VALIDATION LOSS 2.2743, VALIDATION ACCURACY 0.171
[TRAINING EPOCH 6 / 50] TRAIN LOSS 2.2565, TRAIN ACCURACY 0.145 | [VALIDATION EPOCH 6 / 50] VALIDATION LOSS 2.2728, VALIDATION ACCURACY 0.155
[TRAINING EPOCH 7 / 50] TRAIN LOSS 2.2570, TRAIN ACCURACY 0.173 | [VALIDATION EPOCH 7 / 50] VALIDATION LOSS 2.2761, VALIDATION ACCURACY 0.171
[TRAINING EPOCH 8 / 50] TRAIN LOSS 2.2427, TRAIN ACCURACY 0.158 | [VALIDATION EPOCH 8 / 50] VALIDATION LOSS 2.2753, VALIDATION ACCURACY 0.182
[TRAINING EPOCH 9 / 50] TRAIN LOSS 2.2443, TRAIN ACCURACY 0.181 | [VALIDATION EPOCH 9 / 50] VALIDATION LOSS 2.2668, VALIDATION ACCURACY 0.171
[TRAINING EPOCH 10 / 50] TRAIN LOSS 2.2314, TRAIN ACCURACY 0.172 | [VALIDATION EPOCH 10 / 50] VALIDATION LOSS 2.2631, VALIDATION ACCURACY 0.193
[TRAINING EPOCH 11 / 50] TRAIN LOSS 2.1917, TRAIN ACCURACY 0.203 | [VALIDATION EPOCH 11 / 50] VALIDATION LOSS 2.2987, VALIDATION ACCURACY 0.199
[TRAINING EPOCH 12 / 50] TRAIN LOSS 2.1868, TRAIN ACCURACY 0.185 | [VALIDATION EPOCH 12 / 50] VALIDATION LOSS 2.2761, VALIDATION ACCURACY 0.221
[TRAINING EPOCH 13 / 50] TRAIN LOSS 2.1870, TRAIN ACCURACY 0.201 | [VALIDATION EPOCH 13 / 50] VALIDATION LOSS 2.2869, VALIDATION ACCURACY 0.166
[TRAINING EPOCH 14 / 50] TRAIN LOSS 2.1814, TRAIN ACCURACY 0.192 | [VALIDATION EPOCH 14 / 50] VALIDATION LOSS 2.2764, VALIDATION ACCURACY 0.177
[TRAINING EPOCH 15 / 50] TRAIN LOSS 2.1811, TRAIN ACCURACY 0.191 | [VALIDATION EPOCH 15 / 50] VALIDATION LOSS 2.2507, VALIDATION ACCURACY 0.188
[TRAINING EPOCH 16 / 50] TRAIN LOSS 2.1282, TRAIN ACCURACY 0.232 | [VALIDATION EPOCH 16 / 50] VALIDATION LOSS 2.2520, VALIDATION ACCURACY 0.160
[TRAINING EPOCH 17 / 50] TRAIN LOSS 2.1394, TRAIN ACCURACY 0.232 | [VALIDATION EPOCH 17 / 50] VALIDATION LOSS 2.2321, VALIDATION ACCURACY 0.182
[TRAINING EPOCH 18 / 50] TRAIN LOSS 2.1472, TRAIN ACCURACY 0.198 | [VALIDATION EPOCH 18 / 50] VALIDATION LOSS 2.2235, VALIDATION ACCURACY 0.171
[TRAINING EPOCH 19 / 50] TRAIN LOSS 2.1096, TRAIN ACCURACY 0.239 | [VALIDATION EPOCH 19 / 50] VALIDATION LOSS 2.2444, VALIDATION ACCURACY 0.177
[TRAINING EPOCH 20 / 50] TRAIN LOSS 2.0866, TRAIN ACCURACY 0.242 | [VALIDATION EPOCH 20 / 50] VALIDATION LOSS 2.2408, VALIDATION ACCURACY 0.182
[TRAINING EPOCH 21 / 50] TRAIN LOSS 2.0906, TRAIN ACCURACY 0.223 | [VALIDATION EPOCH 21 / 50] VALIDATION LOSS 2.2349, VALIDATION ACCURACY 0.199
[TRAINING EPOCH 22 / 50] TRAIN LOSS 2.0787, TRAIN ACCURACY 0.243 | [VALIDATION EPOCH 22 / 50] VALIDATION LOSS 2.2079, VALIDATION ACCURACY 0.221
[TRAINING EPOCH 23 / 50] TRAIN LOSS 2.0653, TRAIN ACCURACY 0.241 | [VALIDATION EPOCH 23 / 50] VALIDATION LOSS 2.2496, VALIDATION ACCURACY 0.210
[TRAINING EPOCH 24 / 50] TRAIN LOSS 2.0488, TRAIN ACCURACY 0.246 | [VALIDATION EPOCH 24 / 50] VALIDATION LOSS 2.2801, VALIDATION ACCURACY 0.232
[TRAINING EPOCH 25 / 50] TRAIN LOSS 2.0079, TRAIN ACCURACY 0.284 | [VALIDATION EPOCH 25 / 50] VALIDATION LOSS 2.2340, VALIDATION ACCURACY 0.188
[TRAINING EPOCH 26 / 50] TRAIN LOSS 2.0127, TRAIN ACCURACY 0.274 | [VALIDATION EPOCH 26 / 50] VALIDATION LOSS 2.2333, VALIDATION ACCURACY 0.227
[TRAINING EPOCH 27 / 50] TRAIN LOSS 1.9634, TRAIN ACCURACY 0.297 | [VALIDATION EPOCH 27 / 50] VALIDATION LOSS 2.2668, VALIDATION ACCURACY 0.199
[TRAINING EPOCH 28 / 50] TRAIN LOSS 1.9914, TRAIN ACCURACY 0.264 | [VALIDATION EPOCH 28 / 50] VALIDATION LOSS 2.2202, VALIDATION ACCURACY 0.204
```

[TRAINING EPOCH 29 / 50] TRAIN LOSS 2.0037, TRAIN ACCURACY 0.267 | [VALIDATION EPOCH 29 / 50] VALIDATION LOSS 2.2856, VALIDATION ACCURACY 0.204
 [TRAINING EPOCH 30 / 50] TRAIN LOSS 1.9829, TRAIN ACCURACY 0.296 | [VALIDATION EPOCH 30 / 50] VALIDATION LOSS 2.2283, VALIDATION ACCURACY 0.221
 [TRAINING EPOCH 31 / 50] TRAIN LOSS 1.9624, TRAIN ACCURACY 0.289 | [VALIDATION EPOCH 31 / 50] VALIDATION LOSS 2.2457, VALIDATION ACCURACY 0.221
 [TRAINING EPOCH 32 / 50] TRAIN LOSS 1.9550, TRAIN ACCURACY 0.296 | [VALIDATION EPOCH 32 / 50] VALIDATION LOSS 2.2378, VALIDATION ACCURACY 0.232
 [TRAINING EPOCH 33 / 50] TRAIN LOSS 1.9519, TRAIN ACCURACY 0.278 | [VALIDATION EPOCH 33 / 50] VALIDATION LOSS 2.2200, VALIDATION ACCURACY 0.221
 [TRAINING EPOCH 34 / 50] TRAIN LOSS 1.9278, TRAIN ACCURACY 0.306 | [VALIDATION EPOCH 34 / 50] VALIDATION LOSS 2.2262, VALIDATION ACCURACY 0.227
 [TRAINING EPOCH 35 / 50] TRAIN LOSS 1.8877, TRAIN ACCURACY 0.321 | [VALIDATION EPOCH 35 / 50] VALIDATION LOSS 2.2648, VALIDATION ACCURACY 0.232
 [TRAINING EPOCH 36 / 50] TRAIN LOSS 1.9173, TRAIN ACCURACY 0.285 | [VALIDATION EPOCH 36 / 50] VALIDATION LOSS 2.2527, VALIDATION ACCURACY 0.221
 [TRAINING EPOCH 37 / 50] TRAIN LOSS 1.8894, TRAIN ACCURACY 0.311 | [VALIDATION EPOCH 37 / 50] VALIDATION LOSS 2.2410, VALIDATION ACCURACY 0.243
 [TRAINING EPOCH 38 / 50] TRAIN LOSS 1.8857, TRAIN ACCURACY 0.324 | [VALIDATION EPOCH 38 / 50] VALIDATION LOSS 2.2894, VALIDATION ACCURACY 0.227
 [TRAINING EPOCH 39 / 50] TRAIN LOSS 1.8893, TRAIN ACCURACY 0.306 | [VALIDATION EPOCH 39 / 50] VALIDATION LOSS 2.2827, VALIDATION ACCURACY 0.210
 [TRAINING EPOCH 40 / 50] TRAIN LOSS 1.8515, TRAIN ACCURACY 0.315 | [VALIDATION EPOCH 40 / 50] VALIDATION LOSS 2.3031, VALIDATION ACCURACY 0.193
 [TRAINING EPOCH 41 / 50] TRAIN LOSS 1.8276, TRAIN ACCURACY 0.340 | [VALIDATION EPOCH 41 / 50] VALIDATION LOSS 2.2801, VALIDATION ACCURACY 0.232
 [TRAINING EPOCH 42 / 50] TRAIN LOSS 1.8280, TRAIN ACCURACY 0.344 | [VALIDATION EPOCH 42 / 50] VALIDATION LOSS 2.2596, VALIDATION ACCURACY 0.276
 [TRAINING EPOCH 43 / 50] TRAIN LOSS 1.8726, TRAIN ACCURACY 0.349 | [VALIDATION EPOCH 43 / 50] VALIDATION LOSS 2.2377, VALIDATION ACCURACY 0.243
 [TRAINING EPOCH 44 / 50] TRAIN LOSS 1.8222, TRAIN ACCURACY 0.332 | [VALIDATION EPOCH 44 / 50] VALIDATION LOSS 2.2328, VALIDATION ACCURACY 0.243
 [TRAINING EPOCH 45 / 50] TRAIN LOSS 1.8262, TRAIN ACCURACY 0.354 | [VALIDATION EPOCH 45 / 50] VALIDATION LOSS 2.2701, VALIDATION ACCURACY 0.227
 [TRAINING EPOCH 46 / 50] TRAIN LOSS 1.8215, TRAIN ACCURACY 0.332 | [VALIDATION EPOCH 46 / 50] VALIDATION LOSS 2.2242, VALIDATION ACCURACY 0.260
 [TRAINING EPOCH 47 / 50] TRAIN LOSS 1.7908, TRAIN ACCURACY 0.318 | [VALIDATION EPOCH 47 / 50] VALIDATION LOSS 2.3071, VALIDATION ACCURACY 0.243
 [TRAINING EPOCH 48 / 50] TRAIN LOSS 1.7664, TRAIN ACCURACY 0.379 | [VALIDATION EPOCH 48 / 50] VALIDATION LOSS 2.2781, VALIDATION ACCURACY 0.243
 [TRAINING EPOCH 49 / 50] TRAIN LOSS 1.7541, TRAIN ACCURACY 0.369 | [VALIDATION EPOCH 49 / 50] VALIDATION LOSS 2.3296, VALIDATION ACCURACY 0.243
 [TRAINING EPOCH 50 / 50] TRAIN LOSS 1.7784, TRAIN ACCURACY 0.350 | [VALIDATION EPOCH 50 / 50] VALIDATION LOSS 2.2969, VALIDATION ACCURACY 0.249
 TRAINING COMPLETE!!!
 BEST VALIDATION ACCURACY: 0.27624309392265195

In [14]: `plt.figure(figsize=(12, 5))`

```
# Plot 1: Loss Curve
plt.subplot(1, 2, 1)
plt.plot(training_loss, label='Training Loss')
plt.plot(validation_loss, label='Validation Loss')
plt.title('Loss over Epochs')
plt.xlabel('Epochs')
plt.ylabel('Loss')
```

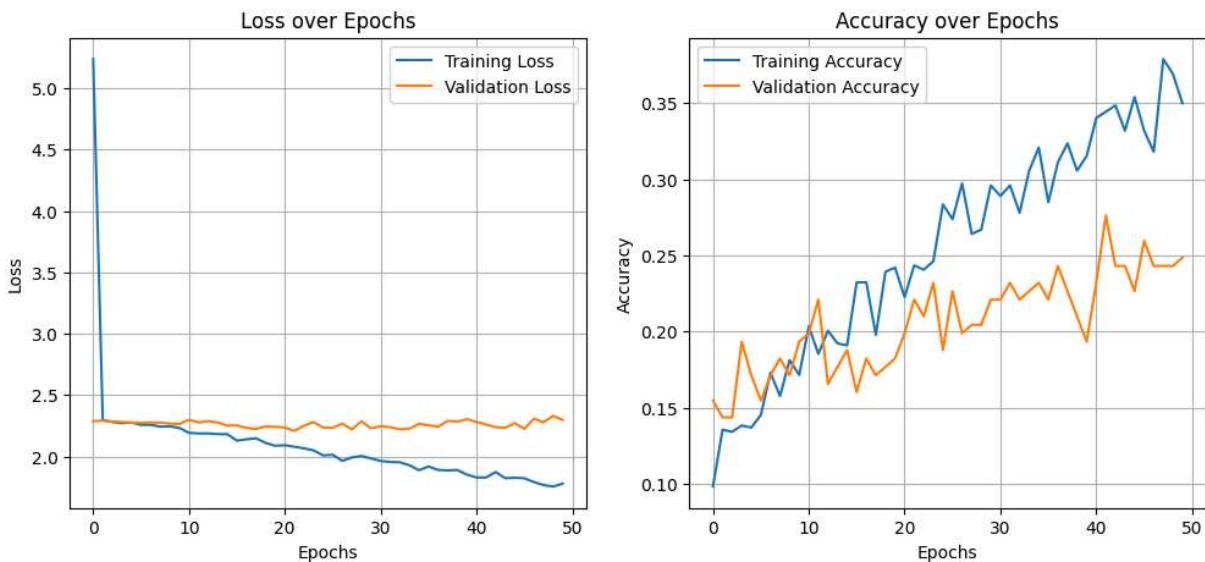
```

plt.legend()
plt.grid(True)

# Plot 2: Accuracy Curve
plt.subplot(1, 2, 2)
plt.plot(training_accuracy, label='Training Accuracy')
plt.plot(validation_accuracy, label='Validation Accuracy')
plt.title('Accuracy over Epochs')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()
plt.grid(True)

plt.show()

```



(7) Make prediction and shows some predicted images

```

In [15]: def display_image(img):
    img = img.cpu().numpy().transpose((1, 2, 0))
    mean = np.array([0.485, 0.456, 0.406])
    std = np.array([0.229, 0.224, 0.225])
    img = std * img + mean
    img = np.clip(img, 0, 1)
    plt.imshow(img)

model.eval()

iterator = iter(validation_loader)

images, labels = next(iterator)

images, labels = images.to(device), labels.to(device)

print(f"Got a batch of {images.size(0)} images.")

```

```

print(f"Image shape: {images.shape}")

with torch.no_grad():
    outputs = model(images)
    _, preds = torch.max(outputs, 1)
fig = plt.figure(figsize=(25, 4))

for idx in range(8):

    ax = fig.add_subplot(1, 10, idx+1, xticks=[], yticks=[])

    display_image(images[idx])

    pred_label = preds[idx].item()
    truth_label = labels[idx].item()

    color = "green" if pred_label == truth_label else "red"

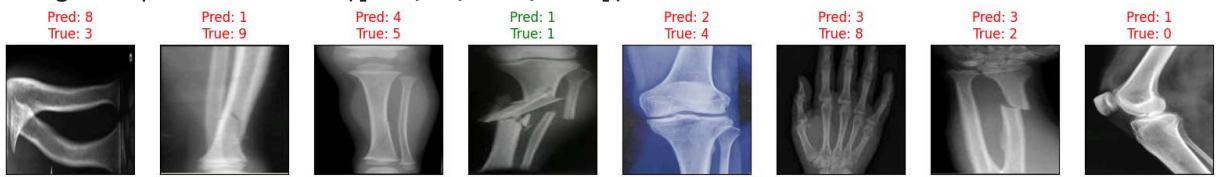
    ax.set_title(f"Pred: {pred_label}\nTrue: {truth_label}", color=color,)

plt.show()

```

Got a batch of 128 images.

Image shape: torch.Size([128, 3, 224, 224])



(8) Make Evaluation with confusion matrix

```

In [16]: all_predictions = []
all_labels = []

model.eval()

with torch.no_grad():
    for images, labels in validation_loader:
        images = images.to(device)
        labels = labels.to(device)

        outputs = model(images)
        _, preds = torch.max(outputs, 1)

        all_predictions.extend(preds.cpu().numpy())
        all_labels.extend(labels.cpu().numpy())

cm = confusion_matrix(all_labels, all_predictions)

plt.figure(figsize=(20, 20))

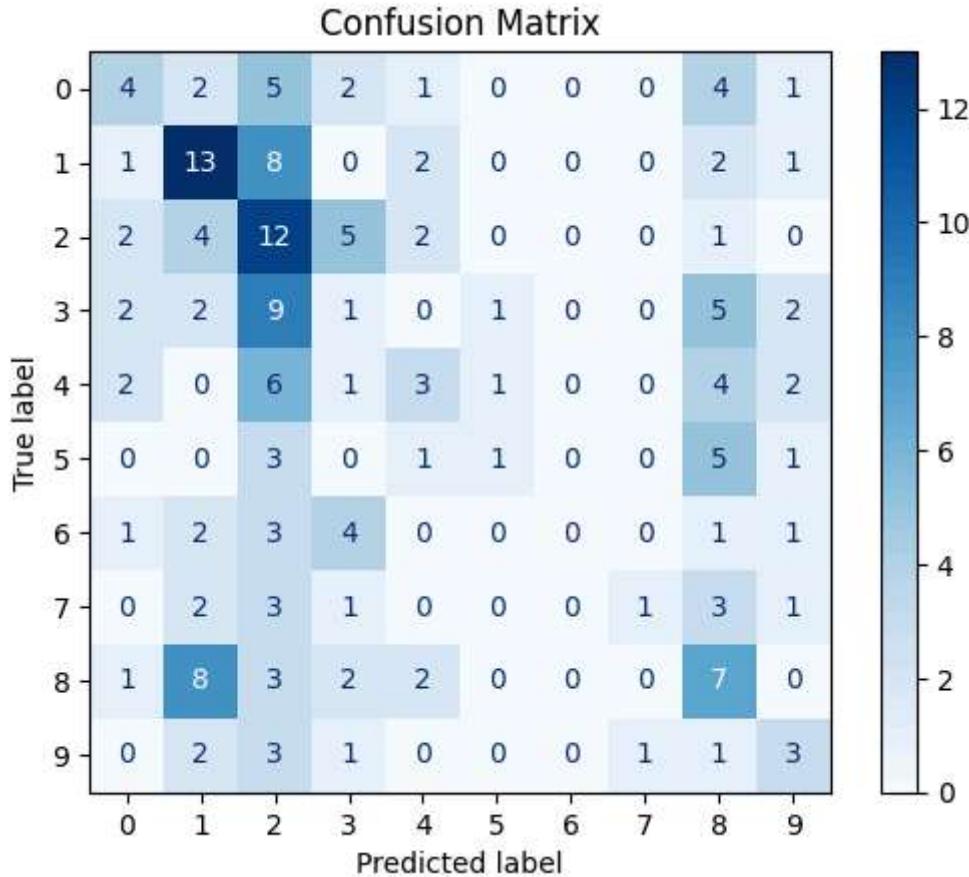
disp = ConfusionMatrixDisplay(confusion_matrix=cm)

```

```
disp.plot(cmap='Blues', values_format='d')

plt.title("Confusion Matrix")
plt.show()
```

<Figure size 2000x2000 with 0 Axes>



(9) Instance of a new CNN

Use the Sequential API it is faster

- This cnn has three hidden layers (conv2d -> proper activation -> max pool 2d) one has 16 units, the second has 32 units, and the third has 64
- Add a flatten layer
- Classifier need 128 Linear unit proper activation and set dropout at 0.5
- The last layer of the classifier how many units need?

Repeat from 5 to 8

```
In [17]: model3 = LayeredCNN([16,32, 64], [50176, 128])

device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

model3.to(device)
```

```

print(f"Model3 is running on: {device}")

criterion = nn.CrossEntropyLoss()

optimizer3 = torch.optim.Adam(model3.parameters(), lr=1e-3, weight_decay=1e-4)

```

Model3 is running on: cuda

```

In [18]: training_loss_3 = []
validation_loss_3 = []
training_accuracy_3 = []
validation_accuracy_3 = []

epochs = 50
best_validation_accuracy_3 = 0.0

for epoch in range(1, epochs + 1):

    train_loss, train_accuracy = train_one_epoch(
        model=model3,
        loader=train_loader,
        criterion=criterion,
        optimizer=optimizer3,
        device=device
    )

    val_loss, val_accuracy = evaluate(
        model=model3,
        loader=validation_loader,
        criterion=criterion,
        device=device
    )

    # Append to the new lists
    training_loss_3.append(train_loss)
    validation_loss_3.append(val_loss)
    training_accuracy_3.append(train_accuracy)
    validation_accuracy_3.append(val_accuracy)

    print(f"[TRAINING EPOCH {epoch} / {epochs}] TRAIN LOSS {train_loss: .4f}, TRAIN ACCURACY {train_accuracy: .4f}, VALIDATION LOSS {val_loss: .4f}, VALIDATION ACCURACY {val_accuracy: .4f}")

    if val_accuracy > best_validation_accuracy_3:
        best_validation_accuracy_3 = val_accuracy
        torch.save(model3.state_dict(), "cnn_3layer_bone_break_classification.pt")

print("TRAINING COMPLETE!!!")
print("BEST VALIDATION ACCURACY (3 Layer): ", best_validation_accuracy_3)

```

```
[TRAINING EPOCH 1 / 50] TRAIN LOSS 2.5866, TRAIN ACCURACY 0.105 | [VALIDATION EPOCH 1 / 50] VALIDATION LOSS 2.2885, VALIDATION ACCURACY 0.133
[TRAINING EPOCH 2 / 50] TRAIN LOSS 2.2786, TRAIN ACCURACY 0.133 | [VALIDATION EPOCH 2 / 50] VALIDATION LOSS 2.2507, VALIDATION ACCURACY 0.138
[TRAINING EPOCH 3 / 50] TRAIN LOSS 2.2752, TRAIN ACCURACY 0.131 | [VALIDATION EPOCH 3 / 50] VALIDATION LOSS 2.2668, VALIDATION ACCURACY 0.171
[TRAINING EPOCH 4 / 50] TRAIN LOSS 2.2464, TRAIN ACCURACY 0.172 | [VALIDATION EPOCH 4 / 50] VALIDATION LOSS 2.2596, VALIDATION ACCURACY 0.144
[TRAINING EPOCH 5 / 50] TRAIN LOSS 2.2311, TRAIN ACCURACY 0.170 | [VALIDATION EPOCH 5 / 50] VALIDATION LOSS 2.2624, VALIDATION ACCURACY 0.160
[TRAINING EPOCH 6 / 50] TRAIN LOSS 2.2274, TRAIN ACCURACY 0.185 | [VALIDATION EPOCH 6 / 50] VALIDATION LOSS 2.2527, VALIDATION ACCURACY 0.155
[TRAINING EPOCH 7 / 50] TRAIN LOSS 2.2363, TRAIN ACCURACY 0.183 | [VALIDATION EPOCH 7 / 50] VALIDATION LOSS 2.2479, VALIDATION ACCURACY 0.166
[TRAINING EPOCH 8 / 50] TRAIN LOSS 2.2039, TRAIN ACCURACY 0.203 | [VALIDATION EPOCH 8 / 50] VALIDATION LOSS 2.2622, VALIDATION ACCURACY 0.138
[TRAINING EPOCH 9 / 50] TRAIN LOSS 2.1876, TRAIN ACCURACY 0.191 | [VALIDATION EPOCH 9 / 50] VALIDATION LOSS 2.2447, VALIDATION ACCURACY 0.155
[TRAINING EPOCH 10 / 50] TRAIN LOSS 2.1779, TRAIN ACCURACY 0.210 | [VALIDATION EPOCH 10 / 50] VALIDATION LOSS 2.2462, VALIDATION ACCURACY 0.133
[TRAINING EPOCH 11 / 50] TRAIN LOSS 2.1505, TRAIN ACCURACY 0.198 | [VALIDATION EPOCH 11 / 50] VALIDATION LOSS 2.2345, VALIDATION ACCURACY 0.149
[TRAINING EPOCH 12 / 50] TRAIN LOSS 2.1540, TRAIN ACCURACY 0.227 | [VALIDATION EPOCH 12 / 50] VALIDATION LOSS 2.2588, VALIDATION ACCURACY 0.144
[TRAINING EPOCH 13 / 50] TRAIN LOSS 2.1562, TRAIN ACCURACY 0.235 | [VALIDATION EPOCH 13 / 50] VALIDATION LOSS 2.2554, VALIDATION ACCURACY 0.144
[TRAINING EPOCH 14 / 50] TRAIN LOSS 2.1264, TRAIN ACCURACY 0.228 | [VALIDATION EPOCH 14 / 50] VALIDATION LOSS 2.2255, VALIDATION ACCURACY 0.144
[TRAINING EPOCH 15 / 50] TRAIN LOSS 2.1217, TRAIN ACCURACY 0.225 | [VALIDATION EPOCH 15 / 50] VALIDATION LOSS 2.2655, VALIDATION ACCURACY 0.122
[TRAINING EPOCH 16 / 50] TRAIN LOSS 2.0846, TRAIN ACCURACY 0.234 | [VALIDATION EPOCH 16 / 50] VALIDATION LOSS 2.2414, VALIDATION ACCURACY 0.149
[TRAINING EPOCH 17 / 50] TRAIN LOSS 2.1032, TRAIN ACCURACY 0.252 | [VALIDATION EPOCH 17 / 50] VALIDATION LOSS 2.2406, VALIDATION ACCURACY 0.155
[TRAINING EPOCH 18 / 50] TRAIN LOSS 2.0739, TRAIN ACCURACY 0.235 | [VALIDATION EPOCH 18 / 50] VALIDATION LOSS 2.2584, VALIDATION ACCURACY 0.149
[TRAINING EPOCH 19 / 50] TRAIN LOSS 2.0618, TRAIN ACCURACY 0.237 | [VALIDATION EPOCH 19 / 50] VALIDATION LOSS 2.2630, VALIDATION ACCURACY 0.127
[TRAINING EPOCH 20 / 50] TRAIN LOSS 2.0513, TRAIN ACCURACY 0.241 | [VALIDATION EPOCH 20 / 50] VALIDATION LOSS 2.2485, VALIDATION ACCURACY 0.160
[TRAINING EPOCH 21 / 50] TRAIN LOSS 2.0276, TRAIN ACCURACY 0.263 | [VALIDATION EPOCH 21 / 50] VALIDATION LOSS 2.2784, VALIDATION ACCURACY 0.193
[TRAINING EPOCH 22 / 50] TRAIN LOSS 1.9963, TRAIN ACCURACY 0.288 | [VALIDATION EPOCH 22 / 50] VALIDATION LOSS 2.2487, VALIDATION ACCURACY 0.177
[TRAINING EPOCH 23 / 50] TRAIN LOSS 1.9946, TRAIN ACCURACY 0.277 | [VALIDATION EPOCH 23 / 50] VALIDATION LOSS 2.2324, VALIDATION ACCURACY 0.188
[TRAINING EPOCH 24 / 50] TRAIN LOSS 1.9654, TRAIN ACCURACY 0.284 | [VALIDATION EPOCH 24 / 50] VALIDATION LOSS 2.2289, VALIDATION ACCURACY 0.204
[TRAINING EPOCH 25 / 50] TRAIN LOSS 2.0170, TRAIN ACCURACY 0.267 | [VALIDATION EPOCH 25 / 50] VALIDATION LOSS 2.1893, VALIDATION ACCURACY 0.188
[TRAINING EPOCH 26 / 50] TRAIN LOSS 1.9539, TRAIN ACCURACY 0.290 | [VALIDATION EPOCH 26 / 50] VALIDATION LOSS 2.2116, VALIDATION ACCURACY 0.177
[TRAINING EPOCH 27 / 50] TRAIN LOSS 1.9071, TRAIN ACCURACY 0.331 | [VALIDATION EPOCH 27 / 50] VALIDATION LOSS 2.2042, VALIDATION ACCURACY 0.193
[TRAINING EPOCH 28 / 50] TRAIN LOSS 1.8910, TRAIN ACCURACY 0.306 | [VALIDATION EPOCH 28 / 50] VALIDATION LOSS 2.2653, VALIDATION ACCURACY 0.215
```

[TRAINING EPOCH 29 / 50] TRAIN LOSS 1.9299, TRAIN ACCURACY 0.315 | [VALIDATION EPOCH 29 / 50] VALIDATION LOSS 2.2013, VALIDATION ACCURACY 0.221
 [TRAINING EPOCH 30 / 50] TRAIN LOSS 1.9048, TRAIN ACCURACY 0.320 | [VALIDATION EPOCH 30 / 50] VALIDATION LOSS 2.2313, VALIDATION ACCURACY 0.232
 [TRAINING EPOCH 31 / 50] TRAIN LOSS 1.8840, TRAIN ACCURACY 0.326 | [VALIDATION EPOCH 31 / 50] VALIDATION LOSS 2.2569, VALIDATION ACCURACY 0.227
 [TRAINING EPOCH 32 / 50] TRAIN LOSS 1.8671, TRAIN ACCURACY 0.335 | [VALIDATION EPOCH 32 / 50] VALIDATION LOSS 2.2673, VALIDATION ACCURACY 0.227
 [TRAINING EPOCH 33 / 50] TRAIN LOSS 1.8499, TRAIN ACCURACY 0.354 | [VALIDATION EPOCH 33 / 50] VALIDATION LOSS 2.2880, VALIDATION ACCURACY 0.227
 [TRAINING EPOCH 34 / 50] TRAIN LOSS 1.8715, TRAIN ACCURACY 0.349 | [VALIDATION EPOCH 34 / 50] VALIDATION LOSS 2.2808, VALIDATION ACCURACY 0.215
 [TRAINING EPOCH 35 / 50] TRAIN LOSS 1.7980, TRAIN ACCURACY 0.371 | [VALIDATION EPOCH 35 / 50] VALIDATION LOSS 2.2546, VALIDATION ACCURACY 0.221
 [TRAINING EPOCH 36 / 50] TRAIN LOSS 1.7620, TRAIN ACCURACY 0.360 | [VALIDATION EPOCH 36 / 50] VALIDATION LOSS 2.2737, VALIDATION ACCURACY 0.204
 [TRAINING EPOCH 37 / 50] TRAIN LOSS 1.7742, TRAIN ACCURACY 0.367 | [VALIDATION EPOCH 37 / 50] VALIDATION LOSS 2.2907, VALIDATION ACCURACY 0.215
 [TRAINING EPOCH 38 / 50] TRAIN LOSS 1.7814, TRAIN ACCURACY 0.353 | [VALIDATION EPOCH 38 / 50] VALIDATION LOSS 2.2866, VALIDATION ACCURACY 0.221
 [TRAINING EPOCH 39 / 50] TRAIN LOSS 1.7398, TRAIN ACCURACY 0.389 | [VALIDATION EPOCH 39 / 50] VALIDATION LOSS 2.2621, VALIDATION ACCURACY 0.238
 [TRAINING EPOCH 40 / 50] TRAIN LOSS 1.7210, TRAIN ACCURACY 0.396 | [VALIDATION EPOCH 40 / 50] VALIDATION LOSS 2.2995, VALIDATION ACCURACY 0.243
 [TRAINING EPOCH 41 / 50] TRAIN LOSS 1.6927, TRAIN ACCURACY 0.394 | [VALIDATION EPOCH 41 / 50] VALIDATION LOSS 2.1934, VALIDATION ACCURACY 0.238
 [TRAINING EPOCH 42 / 50] TRAIN LOSS 1.7052, TRAIN ACCURACY 0.398 | [VALIDATION EPOCH 42 / 50] VALIDATION LOSS 2.2440, VALIDATION ACCURACY 0.238
 [TRAINING EPOCH 43 / 50] TRAIN LOSS 1.6650, TRAIN ACCURACY 0.419 | [VALIDATION EPOCH 43 / 50] VALIDATION LOSS 2.2588, VALIDATION ACCURACY 0.282
 [TRAINING EPOCH 44 / 50] TRAIN LOSS 1.6255, TRAIN ACCURACY 0.432 | [VALIDATION EPOCH 44 / 50] VALIDATION LOSS 2.2642, VALIDATION ACCURACY 0.232
 [TRAINING EPOCH 45 / 50] TRAIN LOSS 1.6381, TRAIN ACCURACY 0.426 | [VALIDATION EPOCH 45 / 50] VALIDATION LOSS 2.2048, VALIDATION ACCURACY 0.254
 [TRAINING EPOCH 46 / 50] TRAIN LOSS 1.5967, TRAIN ACCURACY 0.436 | [VALIDATION EPOCH 46 / 50] VALIDATION LOSS 2.2611, VALIDATION ACCURACY 0.260
 [TRAINING EPOCH 47 / 50] TRAIN LOSS 1.6158, TRAIN ACCURACY 0.433 | [VALIDATION EPOCH 47 / 50] VALIDATION LOSS 2.2436, VALIDATION ACCURACY 0.282
 [TRAINING EPOCH 48 / 50] TRAIN LOSS 1.5833, TRAIN ACCURACY 0.440 | [VALIDATION EPOCH 48 / 50] VALIDATION LOSS 2.2190, VALIDATION ACCURACY 0.215
 [TRAINING EPOCH 49 / 50] TRAIN LOSS 1.5701, TRAIN ACCURACY 0.430 | [VALIDATION EPOCH 49 / 50] VALIDATION LOSS 2.2348, VALIDATION ACCURACY 0.227
 [TRAINING EPOCH 50 / 50] TRAIN LOSS 1.5156, TRAIN ACCURACY 0.452 | [VALIDATION EPOCH 50 / 50] VALIDATION LOSS 2.3087, VALIDATION ACCURACY 0.271
 TRAINING COMPLETE!!!
 BEST VALIDATION ACCURACY (3 Layer): 0.281767955801105

In [19]: `plt.figure(figsize=(12, 5))`

```
# Plot 1: Loss Curve
plt.subplot(1, 2, 1)
plt.plot(training_loss_3, label='Training Loss')
plt.plot(validation_loss_3, label='Validation Loss')
plt.title('Loss over Epochs')
plt.xlabel('Epochs')
plt.ylabel('Loss')
```

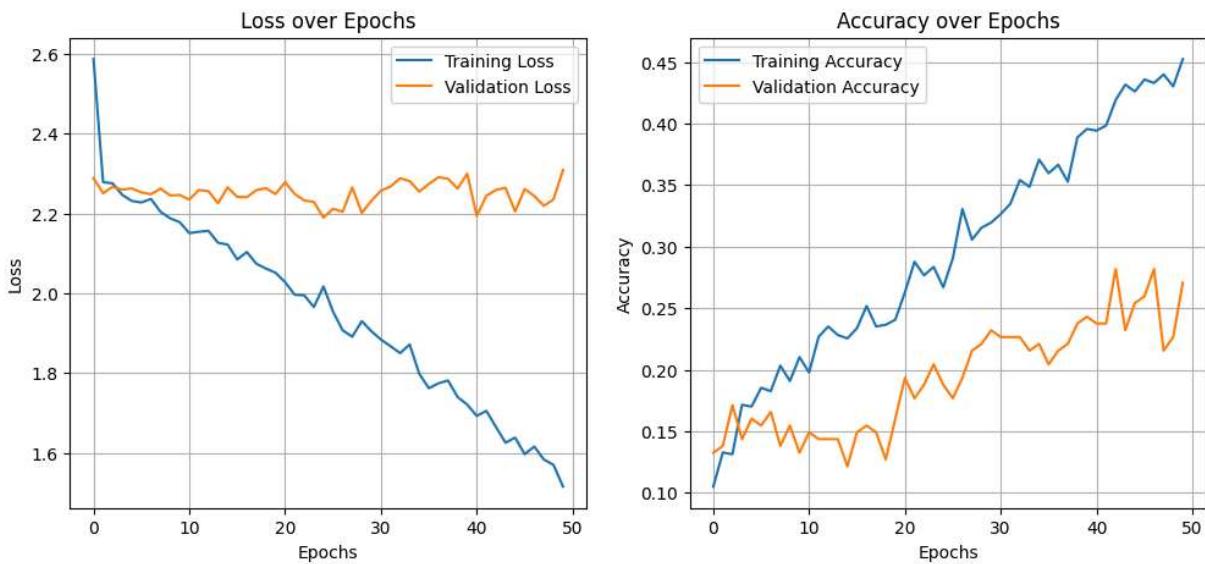
```

plt.legend()
plt.grid(True)

# Plot 2: Accuracy Curve
plt.subplot(1, 2, 2)
plt.plot(training_accuracy_3, label='Training Accuracy')
plt.plot(validation_accuracy_3, label='Validation Accuracy')
plt.title('Accuracy over Epochs')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()
plt.grid(True)

plt.show()

```



```

In [20]: def display_image(img):
    img = img.cpu().numpy().transpose((1, 2, 0))
    mean = np.array([0.485, 0.456, 0.406])
    std = np.array([0.229, 0.224, 0.225])
    img = std * img + mean
    img = np.clip(img, 0, 1)
    plt.imshow(img)

model3.eval()

iterator = iter(validation_loader)

images, labels = next(iterator)

images, labels = images.to(device), labels.to(device)

print(f"Got a batch of {images.size(0)} images.")
print(f"Image shape: {images.shape}")

with torch.no_grad():

    outputs = model3(images)
    _, preds = torch.max(outputs, 1)

```

```

fig = plt.figure(figsize=(25, 4))

for idx in range(8):

    ax = fig.add_subplot(1, 10, idx+1, xticks=[], yticks=[])

    display_image(images[idx])

    pred_label = preds[idx].item()
    truth_label = labels[idx].item()

    color = "green" if pred_label == truth_label else "red"

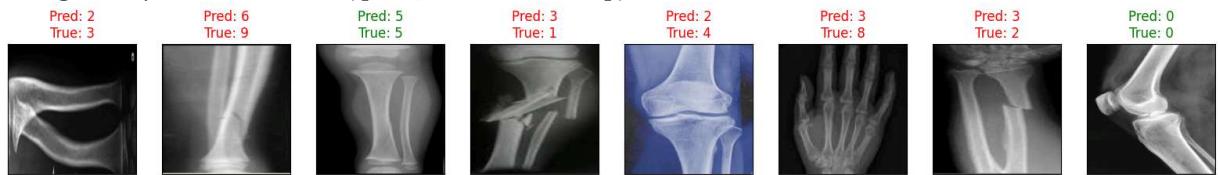
    ax.set_title(f"Pred: {pred_label}\nTrue: {truth_label}", color=color,)

plt.show()

```

Got a batch of 128 images.

Image shape: torch.Size([128, 3, 224, 224])



```

In [21]: all_predictions = []
all_labels = []

model3.eval()

with torch.no_grad():
    for images, labels in validation_loader:
        images = images.to(device)
        labels = labels.to(device)

        outputs = model3(images)
        _, preds = torch.max(outputs, 1)

        all_predictions.extend(preds.cpu().numpy())
        all_labels.extend(labels.cpu().numpy())

cm = confusion_matrix(all_labels, all_predictions)

plt.figure(figsize=(20, 20))

disp = ConfusionMatrixDisplay(confusion_matrix=cm)

disp.plot(cmap='Blues', values_format='d')

plt.title("Confusion Matrix")
plt.show()

```

<Figure size 2000x2000 with 0 Axes>

